

1 When do single-cell response distributions improve drug 2 retrieval?

4 Abstract

5 Artificial intelligence virtual cells aim to predict cellular responses to perturbations and, ultimately,
6 to guide intervention selection. Conventional transcriptomic drug retrieval represents each pertur-
7 bation by a single expression signature, whereas single-cell perturbation atlases preserve cell-to-cell
8 variation in perturbation responses. Here we define population-to-population drug retrieval (PopRe-
9 trieve), in which candidate drugs are ranked by comparing single-cell response populations, and ask
10 when this additional information improves drug ranking. We find that most of the gain over direction-
11 only mean matching is recovered by a mean-based score that also retains response magnitude, with
12 a smaller additional gain from retaining the full response population. The advantage of population-
13 level retrieval weakens as evaluation moves away from response matching: it disappears or reverses
14 under mechanism-of-action recovery, and most of its association with an external functional readout is
15 reproduced by response magnitude alone. The same held-out protein measurements favour different
16 retrieval strategies depending on whether they are summarized as means or retained as populations,
17 showing that the statistical form of an external evaluator can influence which retrieval strategy ap-
18 pears preferable. With candidate responses observed, population-level retrieval corrects an incorrect
19 top-ranked choice about five times as often as it overturns a correct one, but the gain is confined to the
20 roughly one-fifth of queries for which mean matching leaves room for improvement. When candidate
21 responses are predicted rather than observed, this oracle advantage disappears and can reverse sign
22 even when mean-level retrieval performance is preserved. Together, these results show that the value
23 of single-cell population information depends on whether information beyond the mean changes can-
24 didate ordering and, when required, survives forward prediction.

25
26 **Keywords:** artificial intelligence virtual cell; single-cell perturbation; drug retrieval; population-to-
27 population retrieval; cellular heterogeneity; objective–utility mismatch

29 Introduction

30 Artificial intelligence virtual cells^{1,2} aim to predict cellular responses to perturbations^{3–5} and, ultimately,
31 to help prioritize interventions for testing. These goals pose two related problems. Forward prediction
32 asks what response a perturbation will produce^{6–9}; retrieval asks which perturbation best matches a
33 cellular response of interest¹⁰. Transcriptomic drug retrieval has traditionally represented each drug by
34 a single perturbation signature and ranked candidates by similarity or reversal^{11,12}. This approach has
35 made large drug libraries searchable, but collapses variation across cells into a single summary. Single-
36 cell perturbation atlases now capture this variation^{13–15}, including state-dependent responses¹⁶ and ef-
37 fects confined to minority cell populations¹⁷. Here we define population-to-population drug retrieval
38 (PopRetrieve), in which candidate drugs are ranked by comparing single-cell response populations, and
39 ask when this additional information improves drug ranking.

40 More detailed response information does not necessarily lead to a better drug ranking. Two drugs
41 can have similar average responses yet differ in how individual cells respond, creating information that

a population-aware method can use. But an apparent advantage over classical signature matching can arise for more than one reason. Population comparison retains response magnitude that directional similarity discards, so a gain over direction-only matching need not arise from within-population structure. The apparent benefit also depends on how retrieval quality is evaluated: a criterion that preserves full response distributions^{18,19} will naturally favour information that a population-aware scorer retains²⁰⁻²². Such gains establish better response matching, but not necessarily a better biological or functional ranking²³⁻²⁷.

We therefore decompose the problem into information retained, information that changes candidate ordering, and information that survives prediction. We first distinguish three scoring rules by how much of the response each retains: direction, direction plus magnitude, and the full response population. This separates gains recovered by restoring response magnitude from those that require information beyond the mean. We then evaluate the resulting rankings using increasingly task-distal criteria: response matching, mechanism recovery and external functional readouts^{28,29}. We use state-dependent perturbation responses as a concrete mechanism through which population information can affect retrieval. For such information to matter, it must be present and recoverable, change candidate ordering beyond what the mean already determines and, when candidate responses are not observed, survive forward prediction. This separates the existence of single-cell response structure from its value for intervention selection.

Here we show that most of the advantage of population comparison over direction-only signature matching is recovered by a mean-based score that also retains response magnitude, leaving a smaller gain that requires the full response population. The advantage of population comparison weakens as evaluation moves away from response matching: it disappears or reverses under mechanism-of-action recovery, and most of its association with an external functional readout is reproduced by response magnitude alone. The same held-out protein measurements favour different retrieval strategies depending only on whether they are summarized as means or retained as populations, showing that external evaluation is not automatically neutral. In a source-to-target intervention-identification task, population information provides a smaller additional gain beyond magnitude-aware mean matching when candidate responses are observed, and its whole advantage over direction-only mean matching is confined to the queries for which that baseline leaves room for improvement. That advantage does not survive when candidate responses must instead be predicted. Together, these results show that the value of single-cell population information depends not simply on whether additional cellular structure exists, but on whether it changes candidate ordering and survives to the final retrieval decision.

Results

Population-to-population retrieval generalizes classical signature matching

Population-to-population retrieval ranks candidates by comparing query and candidate response populations directly (Fig. 1a,b), whereas classical mean-signature retrieval^{11,12} first collapses each population to a single mean differential-expression vector. Retaining the cells makes response structure available that averaging removes, including differences between cellular states or subpopulations. The retrieval problem is unchanged; what changes is how much of the measured response reaches the ranking.

Two candidate populations can share the same mean response and still differ markedly in how individual cells respond (Fig. 1c,d). A population-aware score separates such candidates where a score computed from the mean cannot, because any such score reads only the mean the two candidates share. We therefore hold the candidate rankings fixed and judge them under criteria that progressively reduce

85 their coupling to the retrieval score: response matching, mechanism recovery and external functional
86 readouts (Fig. 1e,f).

87 Mean-signature and population-level retrieval are one task at two resolutions, connected by the
88 amount of within-population variation retained. Writing a cellular response as

$$X = \mu + \lambda\epsilon, \quad E[\epsilon] = 0,$$

89 the parameter λ controls the variation around the population mean. As λ approaches zero the population
90 collapses onto μ and the energy distance reduces to exactly twice the Euclidean distance between the two
91 population means, a constant that cannot reorder candidates, so the magnitude-aware mean L2 is the
92 zero-variance limit of energy-based population retrieval (Fig. 1g; Methods). The cosine used by classical
93 connectivity scoring discards the length of that mean difference and coarsens the same representation
94 one step further (Fig. 1a,b). The implemented CMap-style cosine applies that same rule to the same
95 mean vectors and is therefore identical to mean-cosine retrieval, whereas the canonical L1000 weighted
96 connectivity score remains a distinct rank-based score (Fig. 2f).

97 The coverage family provides a second continuum: raising β moves the score from averaging sub-
98 population distances toward the worst-matched state, and the single-partition case $K = 1$ recovers the
99 global distance (Fig. 1h). These relationships place classical signature matching and population compar-
100 ison on one axis of sensitivity to within-population structure (Supplementary Table 1).

101 **Most of the gain over directional matching does not require the full response population**

102 Population-level retrieval substantially outperformed direction-only mean matching, and most of that
103 gain did not require the full response population. With candidate responses directly observed, so that
104 the full cellular populations are available to both approaches, energy-distance retrieval reached a Hit@1
105 of 0.836, macro-averaged across the seven task-setting combinations analysed here, against 0.388 for
106 mean-cosine retrieval and the numerically identical CMap-style cosine baseline^{11,12}. Between them sits
107 a magnitude-aware mean L2 at 0.788, applied to the same mean signatures: the step to that score is
108 +0.399 of the gap and the step from it to the full distribution a further +0.048 (Fig. 2a). Because that
109 first step changes both what the score reads, from direction alone to direction and magnitude, and the
110 geometry it uses, from angular similarity to Euclidean distance, mean L2 is a magnitude-aware mean
111 baseline rather than a pure magnitude ablation. The same ordering held in PCA space (Fig. 2a).

112 This ordering was not universal. On Frangieh Perturb-CITE-seq³⁰, the only one of the seven com-
113 binations not built from SciPlex3, the magnitude-aware mean L2 was the best of the nine scores (Hit@1
114 0.611), ahead of mean cosine (0.600) and of energy distance (0.578; Fig. 2b). There, retaining magnitude
115 helped and retaining the full response distribution did not.

116 Population-level scoring also improved response matching under partial observation, where only
117 part of each candidate population is available. The same rankings were judged by a regret criterion that
118 penalizes selecting a candidate whose per-cell outcomes are worse than the best available (Methods).
119 Across 765 queries, the worst-case subpopulation-coverage score reduced median regret by +0.031 rel-
120 ative to direction-only mean cosine (95% bootstrap CI +0.017 to +0.049), with lower regret on 56% of
121 queries (Fig. 2c).

122 Reducing the minority response state from half of the query population to a tenth did not erode the
123 energy-distance advantage consistently: it collapsed in one cell line and held in the other two (Fig. 2d).
124 The pattern belongs to the population-aware family rather than to one score. All five variants improved




Figure 1: Perturbation responses can be represented and evaluated at multiple information resolutions. **a–f**, schematic; **g**, a seeded synthetic construction; **h**, an analytic illustration. **a**, A response population is summarised by its mean μ (orange diamond) or kept as the population P (blue cells); bars, the information each rule retains, in equal-width segments that carry no quantity. **b**, One pair of populations and the three scoring rules: the mean route keeps one vector each, the population route compares the cells. **c**, Under drug B a resistant fifth does not move, yet the two treated means coincide in direction and in magnitude; faint cells, untreated. **d**, Three populations sharing one sample mean exactly, so both mean scores tie by construction while the population score, an energy distance over the drawn cells, prefers A. **e**, Objective-aligned, the score that produced the ranking also grades it; less score-aligned, the same ranking is judged by mechanism of action, target engagement or GDSC2 viability (icons). **f**, Three evaluation levels ordered by coupling to the retrieval score; filled, that level is a study's headline metric. **g**, Energy distance against the mean-only limit as the spread λ is restored, closing at $\lambda = 0$, where the energy distance between two point masses is exactly twice the Euclidean distance between their means. **h**, Four per-subpopulation distances 0.20, 0.55, 0.80 and 1.30, the worst-matched state crossed; frames, the weight the aggregate places on each state at $\beta \rightarrow 0$, $\beta = 2$ and $\beta \rightarrow \infty$.




Figure 2: A magnitude-aware mean score recovers most of the gain beyond directional scoring. Points and intervals are defined in Methods. **a**, Hit@1 for nine scorers, macro-averaged over seven task-setting combinations: direction-only mean cosine 0.388, magnitude-aware mean L2 0.788 and energy distance 0.836, so the two steps are +0.399 and +0.048. Dagger, the published CMap connectivity score (WTCS). **b**, The same three scorers per task ($n = 90, 1,080$ and 90 queries); the magnitude step is +0.52 and +0.41 on the two SciPlex3 tasks and +0.01 on Frangieh, where the distributional step is -0.03. **c**, Worst-case subpopulation coverage against direction-only mean cosine over 765 partial-observation queries; curve, the quantile function of the per-query regret reduction, positive favouring coverage (less regret on 56% of queries, more on 36%, ties on 8%). Median +0.031 (band, 95% bootstrap CI +0.017 to +0.049; paired Wilcoxon $P = 1.6 \times 10^{-12}$, mean +0.132), and +0.035 on the 165 queries with a genuinely incomplete library ($P = 6.0 \times 10^{-5}$). **d**, Energy advantage over mean cosine against the minority fraction α ($n = 20$ seeds per point; top, query composition), falling from +1.00 to 0.00 in K562 and not falling in A549 or MCF7 (+0.40 to +0.65; +0.35 to +0.45). **e**, Median regret reduction for five population-level scorers on 627 diagnostic-positive queries. Against direction-only mean cosine all five 95% bootstrap intervals exclude zero (+0.032 to +0.067); against the magnitude-aware mean L2 the three distributional distances have a median of exactly zero and both coverage scores are negative (-0.030, -0.038). **f**, Spearman correlation among seven scorers over 54,180 query-candidate scores from 1,260 queries, an agreement between scoring rules rather than retrieval evidence; shade, $|\rho|$; rings, the closest pair within a tier. The magnitude-aware mean L2 sits with the distributional scores ($\rho = 0.68$ to 0.80), not with the direction-only scores ($\rho = 0.17$ to 0.20).

on direction-only mean cosine, whereas against the magnitude-aware mean L2 the three distributional distances have a median difference of exactly zero and both coverage scores are negative (Fig. 2e).

A magnitude-aware mean score therefore recovers most of the gain beyond directional scoring, and the further increment from retaining the full population is real and much smaller. Both criteria used so far reward correspondence between response populations, so we next ask whether that increment survives biological and functional evaluation.

Population-level gains weaken under biological and functional evaluation

Whether that additional information improves the same candidate rankings depends on how success is defined. We therefore kept the rankings fixed and changed only the criterion used to judge them (Fig. 3).

Mechanism-of-action recovery is the first task-proximal biological criterion. Among the 600 partial-observation queries carrying both response-matching and mechanism annotations, the worst-case subpopulation-coverage scorer improved on direction-only mean cosine by a median of +0.028 under response matching, whereas the identical rankings judged by mechanism-of-action recovery gave a median of 0.000 and a mean of -0.024 (Fig. 3a,b). At the query level, mechanism recovery decreased more often than it increased (40.8% against 35.7%). The same divergence between the two criteria appeared in all five population-level scoring variants, whose response-matching median gains span +0.028 to +0.067 while their mechanism-recovery means span -0.011 to -0.033. Gains under response matching therefore did not consistently carry over to the recovery of biologically related drugs.

A second task-proximal criterion tests how well population-level retrieval preserves minority-state responses. Across 239 task instances, population-level retrieval improved minority-state coverage by only +0.0011 on average, and in 102 of them the two approaches scored exactly the same (Fig. 3c). The small positive effect was concentrated in the constructed cross-cell-line tasks (+0.0037), whereas the remaining settings were near zero or slightly negative. Stratifying by measured response divergence revealed a gradient only on the objective-aligned criterion: the minority-state advantage rose from below zero in the least divergent quartile to about +0.005 in the two most divergent (Fig. 3d), while no divergence stratum reached a mechanism-of-action gain whose interval lay above zero (Fig. 3e). Stronger population differences therefore do not by themselves produce a more biologically useful ranking, and the pre-specified diagnostic gate did not identify the queries that did benefit (Fig. 3f,g; Supplementary Note 1).

Population-level retrieval does carry information that an external functional readout can see. Using GDSC2^{28,29} dose-response profiles from cell lines excluded from the transcriptional retrieval experiment, and removing the shared cell-line sensitivity effect, energy-distance retrieval reached a Spearman correlation of +0.265 with drug-drug functional similarity, against +0.083 for the corresponding mean-signature baseline (Fig. 3h; Supplementary Table 3).

Much of that advantage did not require comparing cellular distributions. Ranking candidates only by how closely their overall response magnitude matched the query reached +0.232 from a single scalar per perturbation, and once that magnitude channel was accounted for the remaining association for energy retrieval was +0.105, positive in all three cell lines (Fig. 3h,i). The channel is directly measurable: across queries the query-candidate energy distance tracks the candidate's own response magnitude at a median Spearman correlation of +0.780 (Fig. 3k), and scored against absolute GDSC2 potency instead, energy and raw mean cosine turn negative, since matching a weak query's magnitude returns other weak drugs, whereas ranking by candidate magnitude alone turns positive (Fig. 3j). Population-level retrieval therefore carries external functional signal, and most of the signal detected here is recoverable




Figure 3: The population advantage weakens, and can change direction, once the judge stops sharing the retrieval objective. Points, medians; bars, 95% bootstrap intervals unless stated; cell lines, line style in **b**, marker shape in **i**. **a**, 600 leave-one-drug-out queries judged twice: median +0.028 under response matching against 0.000 under mechanism-of-action recovery (24% ties; Wilcoxon $p = 7.7 \times 10^{-3}$). **b**, Per cell line, medians exactly 0.000 in all three ($n = 200$ each). **c**, Minority-state coverage gain over 239 tasks, 102 of which score identically; diamonds, dataset means; daggers, constructed data. **d**, Population-minus-mean gain by divergence quartile, objective-aligned judge ($n = 765$): -0.0035 (CI -0.0076 to +0.0001), +0.0038, +0.0054, +0.0051 (Spearman $\rho = +0.14$, $p = 1 \times 10^{-4}$). **e**, The same quartiles under mechanism-of-action recovery ($n = 600$): no interval lies above zero, the two least divergent wholly below it (-0.050, -0.039; $q = 0.002, 0.025$, Benjamini-Hochberg). **f**, Regret reduction, 627 recommended queries against 127 declined (-0.008, CI -0.032 to +0.056, Mann-Whitney $p = 0.64$). **g**, Below, structure reliability against divergence ($\rho = -0.19$, $P = 1.5 \times 10^{-7}$); above, declined queries are more divergent (1.68 against 1.66, $P = 0.009$). **h**, Ranking against GDSC2 drug-drug similarity ($n = 103$); filled, ranking methods; arc, energy partialled on magnitude. **i**, Per query: energy beats the magnitude scalar on 62 queries and loses on 40 (paired $p = 0.054$). **j**, Rankings against absolute GDSC2 potency: query-magnitude matching turns negative, candidate magnitude alone positive. **k**, Per-query Spearman between energy distance and candidate response magnitude, median +0.780.

168 without any cell-to-cell distributional structure.

169 The apparent value of population-level retrieval therefore depends on how success is measured, and
170 it does so even when the criterion moves beyond transcriptomic response matching.

171 **Evaluator form can reverse the preferred retrieval strategy, and natural heterogeneity leaves little** 172 **reordering headroom**

173 We next asked whether the dependence on the evaluation criterion persists when rankings are judged
174 by a biological measurement that no retrieval method can see.

175 Frangieh Perturb-CITE-seq³⁰ measures RNA and 20 surface proteins in the same 218,331 cells, so an
176 RNA-derived ranking can be judged by a modality withheld from every retrieval method. We built two
177 protein-based evaluators from those same measurements, changing only how they were summarized,
178 and scored the identical two RNA-derived rankings against both (Fig. 4a). The first compared the mean
179 protein responses of each query–candidate pair, and under it mean-signature retrieval outperformed
180 energy-distance retrieval (+0.242 versus +0.222).

181 Against the full distributions of the same protein responses the ordering reversed: energy-distance
182 retrieval reached +0.393 against +0.230 for mean-signature retrieval (Fig. 4b). Energy led under the
183 distribution-based evaluator in all three immune conditions, and the change in which method wins
184 occurred in two of the three. Changing only how the same held-out biological measurements were
185 represented was therefore enough to reverse which retrieval strategy performed better.

186 Part of that reversal needs no distributional comparison. A scalar score that compares only the over-
187 all magnitude of the query and candidate responses rose from +0.125 under the mean-based protein
188 evaluator to +0.276 under the distribution-based one. Energy distance stayed ahead of that control un-
189 der both forms, by +0.098 and +0.117 and in every immune condition separately (Fig. 4c). Response
190 magnitude therefore explains part of the reversal, while energy retains an additional association beyond
191 that scalar channel. The same held-out measurements can favour different retrieval strategies depending
192 on how their response structure is represented.

193 A controlled HIR-Bench construction confirms that subpopulation structure can change the preferred
194 candidate where state-specific preferences conflict, although its analytical boundary did not transfer to
195 the biological datasets (Fig. 4d; Supplementary Note 2).

196 We next asked whether state-dependent perturbation responses are recoverable in naturally hetero-
197 geneous tissue¹⁶. In glioblastoma samples from ZhaoSims2021, the malignant and myeloid compart-
198 ments of each tumour were analysed separately against compartment-matched controls. A supervised
199 probe recovered within-patient drug-response identity at an accuracy of 0.923 against 0.777 for the best
200 unsupervised procedure, a median paired gap of +0.117 against +0.007 in a constructed comparator
201 (Fig. 4e; Supplementary Note 3).

202 Whether that structure is large enough to change which drug should be preferred is a separate ques-
203 tion, and the tumour data constrain it. Across 18 within-patient drug pairs, myeloid-response similarity
204 ranked malignant-response similarity at Spearman $\rho = +0.835$ across the two disjoint compartments,
205 which share no cells (Fig. 4f). Drug-response similarity was therefore largely preserved across distinct
206 cellular compartments, leaving less room for state-specific structure to reorder candidates than the pres-
207 ence of heterogeneity alone would suggest. This measures preservation of drug ordering rather than
208 retrieval performance, and it is the requirement the intervention-identification task below tests directly:
209 cellular structure must not only exist and be recoverable, but must also change candidate preference.




Figure 4: External evaluation is not automatically neutral, and natural heterogeneity does not automatically imply decision reordering. Full definitions and robustness analyses are in Supplementary Notes 2 and 3. **a–c**, One external modality judging the same rankings two ways. **a**, One Perturb-CITE-seq experiment measures RNA and 20 surface proteins in the same 218,331 cells; the two RNA rankings and the control-centred protein responses are held identical while only the statistical form of the protein evaluator changes. **b**, Spearman ρ of three RNA-side scores against that evaluator, built once as a cosine between mean protein deltas and once as an energy distance between protein cell clouds; each bar is the median over $n = 659$ queries of a per-query rank correlation across that query's candidate knockouts. Open bar, the response-magnitude control, which ranks candidates but compares no distributions. **c**, How far energy exceeds the magnitude control under each form, pooled and within each immune condition (ratios 1.06 to 1.84); no uncertainty is shown, because the source records one correlation per condition per form. **d**, The HIR-Bench boundary $\alpha^* = B/(A+B)$ separates settings in which the mean is sufficient from settings in which subpopulation structure can change the preferred candidate, without identifying them in biological data (Supplementary Note 2). **e, f**, The same question in four glioblastoma tumours, each split by cell identity into disjoint malignant and myeloid compartments that share no cells and see the same drug and matched control. **e**, Response-identity recoverability of within-patient drug-response structure; open bars, best unsupervised method; filled bars, supervised probe; median paired gap +0.117 (95% patient-level cluster bootstrap +0.115 to +0.135, $n = 36$ splits), against +0.007 in a constructed comparator. **f**, Myeloid-response similarity against malignant-response similarity across those disjoint compartments, on which no retrieval was performed (Spearman $\rho = +0.835$, $n = 18$ drug pairs; dashed line, equal similarity).

210 Population information changes an intervention decision when candidate responses are observed

211 A retrieval system that is to propose interventions must rank compounds whose effect in the query con-
212 text has not been measured. We therefore ran the same comparison as a source-to-target intervention-
213 identification task on Tahoe-100M¹⁵: a held-out cell line’s untreated cells and a desired treated popula-
214 tion enter, and a ranking of a fixed 92-compound library leaves, with the true intervention’s rank as the
215 outcome (3,992 queries over 44 cell lines; Fig. 5a). Because the target is an observed response produced
216 by one library compound, the task measures intervention identification rather than therapeutic utility.
217 It was run twice, once with candidate responses observed within the query line and once with them
218 predicted from the other 43 lines, which separates the value of population information from the cost of
219 forecasting it.

220 With candidate responses observed, the three scoring rules of Fig. 1a again form a ladder. Mean
221 reciprocal rank was 0.945 for direction-only mean cosine, 0.965 once response magnitude was restored
222 and 0.976 for the energy distance between populations, so the step to the magnitude-aware mean is
223 +0.0201 and the step from it to the population +0.0102 (Fig. 5b). The advantage reaches individual
224 decisions: against the magnitude-aware mean, population scoring put the true drug first on 2.0% of
225 queries where the mean route did not and broke a correct choice on 0.5%, a fourfold asymmetry, and
226 against direction-only matching the two rates are 5.6% and 1.1% (Fig. 5c).

227 The pooled oracle gain is limited by the headroom that mean matching leaves. On 3,110 of the
228 3,992 queries the direction-only mean route already ranks the true drug first at every seed, where the
229 population route can only lose and does, by -0.0036 on average; the remaining 882 queries more than
230 account for the pooled $+0.0303$, gaining $+0.1501$ (Fig. 5d). Headroom of that kind correlates with the
231 population route’s gain at Spearman $\rho = +0.79$, but a permutation null that preserves the ranking
232 ceiling reproduces it at $+0.786$ ($+0.767$ to $+0.805$; Fig. 5e, Methods), so the correlation measures the
233 room available for improvement rather than which queries benefit.

234 We next asked which measurable query properties explain the remaining advantage. A query drug’s
235 own treated cells are separable from vehicle cells of the same line at a median accuracy of 0.800 when the
236 labels are given and 0.555 without them, on a balanced two-class problem where 0.5 is chance (Fig. 5f),
237 and drug-by-state interaction was measured by a cross-fitted statistic that is zero in expectation under
238 an additive null and largely free of the cell-count and magnitude dependence a naive response cosine
239 carries (Fig. 5g; Methods). Conditional on the mean route’s own reciprocal rank, response detectability
240 is the only one of the three whose interval excludes zero ($+0.0091$, 95% CI $+0.0060$ to $+0.0121$, $n = 3,618$
241 queries in 44 cell-line clusters); neither the interaction share ($+0.0015$, $p = 0.21$) nor its reproducibility
242 across disjoint halves (-0.0011 , $p = 0.21$) is associated with the population route. Substituting the
243 naive cosine for the cross-fitted statistic reverses the sign of the association (-0.0096 , 95% CI -0.0167 to
244 -0.0024 ; Fig. 5h). Population-level information therefore has measurable value at the oracle, and little
245 room in which to act.

246 The oracle advantage does not survive forward prediction

247 Replacing the observed candidate responses with predicted ones removes the advantage and reverses its
248 sign (Fig. 6a). The primary average-effect predictor costs the mean route 0.058 MRR and the population
249 route 0.124, and the two cross (Fig. 6b). Under it the population route sits -0.0356 below direction-
250 only mean matching (95% CI -0.0620 to -0.0115), so the oracle gain does not attenuate but changes
251 sign (Fig. 6c). None of the six forward predictors evaluated delivers a population advantage, and the
252 Gaussian optimal-transport map is the only one whose point estimate is not negative ($+0.0027$, interval




Figure 5: Population structure adds a small oracle retrieval gain beyond response magnitude, and most of it comes from a fifth of the queries. Source-to-target intervention identification on Tahoe-100M: 3,992 queries over 44 cell lines and a fixed 92-compound library, with candidate responses *observed* within the query line throughout and predicted from the other 43 lines in Fig. 6; **a** is schematic and energy distances use the unbiased U-statistic estimator. **a**, A held-out line's untreated cells and a desired treated population enter and a ranking of all 92 compounds leaves; the target population is the query, not a candidate response. **b**, Mean reciprocal rank for the three scoring rules of Fig. 1a, with 95% bootstrap intervals: 0.945, 0.965 and 0.976, so the two steps are +0.0201 and +0.0102. **c**, Top-1 decisions changed against each mean reference, ratios above the pairs; corrected, the population route puts the true drug first where the reference did not, and broken is the reverse. **d**, The same +0.0303 pooled advantage split by whether the mean route left any room; on 3,110 of 3,992 queries it already ranks the true drug first at every seed. **e**, Spearman correlation between that headroom and the population gain, against a null that reshuffles the population route's reciprocal ranks within each context, keeping both marginals, the $RR \leq 1$ ceiling and the clustering. **f**, Single-cell response detectability, separating a query drug's own treated cells from vehicle cells of the same line at up to 300 cells per arm; median accuracy 0.800 with the labels given and 0.555 without them, on a balanced two-class problem where 0.5 is chance. **g**, Each interaction statistic against the two quantities it must not measure: the naive $1 - \cos(r_1, r_2)$ runs at -0.69 with response magnitude and -0.43 with cells per arm, the cross-fitted S_{int} at +0.66 and +0.10. **h**, Standardised coefficients from regressions of the population route's own reciprocal rank on each explainer, with the mean route's reciprocal rank as a covariate; the first three rows are one fit ($n = 3,618$ queries in 44 cell-line clusters) and the naive cosine a separate one ($n = 3,620$). Filled markers, coefficients whose interval excludes zero.




Figure 6: The oracle advantage does not survive forward prediction, and the deficit is a magnitude error rather than a distributional one. The same task, queries, library and split as Fig. 5, with candidate responses in the held-out line produced by a forward predictor fitted on the other 43 lines; panel **a** is schematic and intervals are 95% bootstrap intervals over cell-line clusters. **a**, The task with that one station replaced; every other element is drawn as in Fig. 5a. **b**, Both routes travelling from observed to predicted candidates under the primary average-effect predictor; the mean route gives up 0.058 MRR and the population route 0.124, and the two cross. **c**, Δ MRR of the population route over direction-only mean matching, at the oracle and for six predictors; under an additive predictor the oracle gain inverts rather than attenuating. **d**, The same predictions judged three ways: the population route's -0.0356 deficit against direction-only is matched by the magnitude-aware mean's -0.0335 , and -0.0020 remains once magnitude is controlled. **e**, Median drug-by-state interaction in each predictor's own output against the cross-fitted value measured in real cells (dashed); an additive predictor's is zero to float precision, an algebraic property of the model class. **f**, A state-conditioned average-effect predictor is the first to generate real interaction without giving up mean-route accuracy (panel **e**), and it does not close the gap. **g**, The same attrition counted over the 755 queries on which the oracle population route had any gain: 209 keep a gain under the primary average-effect predictor and 546 do not. Gains also appear where the oracle had none, 343 for that predictor and 1,165 for the Gaussian OT map, on a different denominator and, for that map, on a mean route reaching only 0.665. **h**, The budget across both figures; the first two rows share the oracle MRR unit and are drawn to scale, and the last two are a rate and a signed gain and carry no width.

dictions the magnitude-aware mean loses -0.0335 to direction-only matching, almost the whole of the population route's deficit, although it reads no distribution at all; against the magnitude-aware mean itself the population route sits only -0.0020 below (-0.0036 to -0.0005 ; Fig. 6d). An additive predictor translates the query line's own vehicle cells by an effect averaged over the other 43 lines, so it reproduces that population's heterogeneity exactly while generating no drug-by-state interaction at all, an algebraic property of the model class rather than a training failure (Fig. 6e; Supplementary Note 4).

A state-conditioned average-effect predictor is the first to generate real interaction while preserving the mean route's MRR (0.8883 against 0.8876), and it still leaves the population route behind (-0.0314 , -0.0584 to -0.0076 ; Fig. 6f). Generating state-specific interaction is therefore necessary for the population route but not sufficient. The same attrition is visible query by query: of the 755 queries on which the oracle population route had any gain, the average-effect predictor keeps one on 209 and loses it on 546 (Fig. 6g).

Together the two arms locate where population-level information enters and where it is lost (Fig. 6h). With responses observed it is real, it reaches a ranking decision, and it is bounded by how good mean signatures already are. With responses predicted it does not survive: the predictors that best preserved mean-level retrieval were additive and generated no state-specific perturbation response, and introducing state conditioning restored that interaction without restoring a population-level advantage.

Discussion

The results trace one chain of attrition rather than a checklist. Response populations carry retrievable information; a smaller part of it exceeds response magnitude; a smaller part again changes which candidate is preferred; and almost none of that survives once candidate responses must be predicted rather than observed. The binding constraint changes along that chain: at the oracle it is the ranking headroom that mean matching leaves, because mean matching already ranks the true compound first on most queries, and once candidate outcomes are unobserved it is the fidelity of the forward predictor.

Single-cell perturbation atlases¹³⁻¹⁵ change the representation available for drug retrieval: a perturbation can now be described by a population of cellular responses rather than by a single averaged signature. Across the datasets and evaluators examined here, retaining the full population increased the information available to the retriever, but only a smaller component beyond response magnitude consistently reached the ranking. The central distinction is therefore not mean versus distributional retrieval, but information that is available versus information that changes the decision. Classical perturbation signatures and single-cell populations are points on a continuum of retained response structure, and the relevant question is how much of that structure the decision at hand requires.

This distinction has direct implications for artificial intelligence virtual cells^{1,2}. A virtual cell that accurately predicts an average perturbation response can still fail to provide the information needed for population-level intervention selection. A model that applies the same perturbation effect to every cell can preserve heterogeneity in the starting population yet, by construction, cannot generate state-specific differential responses; the limit is exact for strictly additive predictors, although more flexible latent models are not algebraically constrained to it (Supplementary Note 4). A downstream population-aware retriever cannot recover variation that the predictor never generated. Reproducing state-dependent differential response is therefore necessary to exploit that particular source of population information, and it is not sufficient: introducing state conditioning restored measurable drug-by-state interaction while preserving mean-level performance, yet did not restore a population-level retrieval advantage (Fig. 6e,f). The same experiment identifies where to look next, because the magnitude-aware mean loses almost ex-

298 actly what the population route loses, so the dominant remaining error is predicted response magnitude
299 rather than population shape (Fig. 6d). That is a concrete target for virtual-cell modelling.

300 Two limits bound these conclusions. The response and interaction diagnostics examined here did
301 not yield a reliable query-level rule for when population scoring helps, and state-dependent differential
302 response is only one of the forms population structure could take, so a gain arising from another form is
303 not excluded. These results also suggest a practical principle for evaluating population-aware retrieval.
304 The statistical representation of an external readout can itself change which method ranks higher^{27,31},
305 so validation should ask whether the added cellular information changes an outcome relevant to the
306 intended use, and should compare against simpler channels such as response magnitude²³ before at-
307 tributing a gain to population structure itself. The functional readouts analysed here remain indirect
308 and do not establish therapeutic benefit. More broadly, a richer cellular measurement justifies a richer
309 retrieval method only when the information it adds survives to the decision that method exists to make.

310 **Methods**

311 **Study design and data resources**

312 This study asks when retaining the full population of single-cell perturbation responses improves re-
313 trieval beyond a conventional mean expression signature. The analyses were therefore organized around
314 four distinct roles: observed-population retrieval, evaluation against biological and functional criteria,
315 measurement of naturally occurring response heterogeneity, and large-scale assessment of the condi-
316 tions under which population-level information can be useful. These roles were analysed separately
317 rather than pooled into a single benchmark.

318 The primary observed-population retrieval analyses used SciPlex3¹⁴, comprising 276,325 cells across
319 three cancer cell lines (A549, K562 and MCF7), 188 compounds shared across the three contexts and
320 a 2,000-gene expression space. SciPlex3 supports both within-context retrieval and controlled cross-
321 context constructions because the same compound library is measured across all three cell lines. Frang-
322 ieh melanoma Perturb-CITE-seq³⁰ spans Control, IFN γ and Co-culture conditions with 239 CRISPR
323 knockouts. Both analyses read the same RNA matrix under two filters: the retrieval analysis uses the
324 218,023 cells left after dropping nine knockout targets with fewer than 50 cells, and the evaluator-shape
325 experiment uses all 218,331 cells of the RNA-protein barcode intersection, of which 215,915 contribute
326 to at least one condition mean once the 60-cell floor per knockout is applied. We use this dataset as
327 an independent perturbation-retrieval setting and, because matched surface-protein measurements are
328 available, for the evaluator-shape experiment described below, in which that protein modality is held
329 out from every retrieval method and used only to judge. It is a genetic-perturbation test of the same
330 retrieval formulation rather than a drug-retrieval benchmark. CD34+ haematopoietic progenitor data
331 (GSE306429)³² contain 33,984 cells and 36 compounds and are used as a negative-control setting for
332 population-level retrieval.

333 Three additional resources serve roles outside the primary retrieval benchmark. ZhaoSims2021¹⁶
334 contains drug-treated patient-derived glioblastoma tissue and is used to test whether state-dependent
335 responses and their recoverability are present in naturally heterogeneous tissue; no retrieval perfor-
336 mance is reported from this cohort. Tahoe-100M¹⁵ serves two distinct roles: it provides larger-scale
337 measurements of the same information requirements under intrinsic within-cell-line heterogeneity, and
338 it is the substrate for the source-to-target intervention-identification benchmark of Figs. 5 and 6. We
339 analyse plate 3 only, which comprises 4,158,278 cells across 50 cell lines and 93 compounds, using the
340 released 2,304-gene log1p-transformed representation without refitting the preprocessing pipeline. Fi-

nally, GDSC2 dose-response profiles provide an external functional readout for drug-drug similarity and are joined to the SciPlex3 compound panel independently of the transcriptomic retrieval scores.

Observed responses, predicted responses, external functional measurements and natural-tissue analyses therefore answer different questions in this study. Results are not combined across these resources as though they represented independent replicates of one benchmark. Dataset scale, inclusion criteria, experimental role and query construction are summarized in Supplementary Table 2.

Preprocessing and construction of perturbation-response populations

Single-cell expression analyses were performed in log-transformed expression space using the processed gene sets defined for each dataset. SciPlex3, Frangieh and CD34+ were analysed in 2,000-gene representations. For SciPlex3 and Frangieh these genes were selected in this study rather than inherited from the released objects: counts were normalized to a total of 10^4 per cell and log1p-transformed, 2,000 highly variable genes were chosen by the Seurat criterion³³ on a subsample of cells, and the final matrices were then rebuilt from raw counts on those genes using all-gene size factors so that normalization is applied once. SciPlex3 was additionally downsampled to at most 120 cells per compound, cell line and dose, and 3,000 control cells per cell line. The CD34+ tensor was received in processed form and the feature-selection procedure applied to it is not documented by the source. For ZhaoSims2021, counts were normalized to a total of 10^4 per cell, log1p-transformed and reduced to 2,000 highly variable genes. Tahoe-100M was used in the released preprocessed form (2,304 highly variable genes, log1p transformed), with no additional feature selection or representation fitting. Dataset-specific filtering and minimum-cell requirements for individual analyses are given in the corresponding Methods sections below.

For a perturbation d measured in cellular context c , let

$$X_{d,c} = \{x_1^{(d,c)}, \dots, x_{n_{d,c}}^{(d,c)}\}$$

denote its treated single-cell expression profiles and let

$$\mu_{0,c} = \frac{1}{n_{0,c}} \sum_{j=1}^{n_{0,c}} x_j^{(0,c)}$$

denote the mean of the matched control population. The conventional perturbation signature is the control-referenced mean response

$$\bar{d}_{d,c} = \frac{1}{n_{d,c}} \sum_{i=1}^{n_{d,c}} x_i^{(d,c)} - \mu_{0,c}.$$

Mean-signature retrieval operates only on $\bar{d}_{d,c}$.

Population-level retrieval retains the corresponding control-referenced cellular responses,

$$P_{d,c} = \left\{ x_i^{(d,c)} - \mu_{0,c} \right\}_{i=1}^{n_{d,c}},$$

so that the mean of $P_{d,c}$ is exactly the conventional mean perturbation signature while variation among its cells remains available to the retrieval score. This construction makes the comparison between mean-signature and population-level retrieval differ in the response structure retained rather than in the underlying perturbation or control definition.

Controls were matched to the biological context in which the perturbation was measured. For anal-

yses that explicitly compare responses between cellular states or compartments, each state is instead referenced to its own matched control population so that pre-existing differences between states cannot be mistaken for differential perturbation responses. This state-specific construction is used for the differential-response analyses in predicted populations, patient-derived glioblastoma and Tahoe-100M and is defined in the corresponding sections below.

Cell sampling was performed only after the response populations had been defined. Minimum cell numbers, per-condition caps and random seeds differ by experiment and are reported with the corresponding analysis rather than imposed globally. Protein measurements used for the evaluator-shape experiment were processed separately, including barcode-based RNA-protein matching and protein-specific normalization, as described below.

Population-to-population retrieval and signature baselines

For a query response population Q and a library of candidate response populations $\{P_d\}_{d=1}^D$, retrieval assigns each candidate a similarity or distance to the query and orders the candidate library accordingly. Population-level methods operate directly on the empirical response populations defined above, whereas signature-based methods first reduce each population to its control-referenced mean response. Query construction and candidate libraries are kept identical when comparing the two representations.

Population-level retrieval. We evaluated energy distance³⁴ (the primary population score), maximum mean discrepancy³⁵, sliced Wasserstein distance³⁶, and the mean and worst-case limits of a subpopulation-coverage score. All rank candidates from smallest to largest distance. For a candidate population $P = \{x_1, \dots, x_m\}$ and a query population $Q = \{y_1, \dots, y_n\}$ the energy distance is estimated as

$$\hat{\mathcal{E}}(P, Q) = \frac{2}{mn} \sum_{i,j} \|x_i - y_j\| - \frac{1}{m(m-1)} \sum_{i \neq i'} \|x_i - x_{i'}\| - \frac{1}{n(n-1)} \sum_{j \neq j'} \|y_j - y_{j'}\|,$$

the Euclidean norm itself rather than its square. Excluding the self-pairs makes this the unbiased U-statistic form; the biased V-statistic divides the two within-population terms by m^2 and n^2 instead, which adds each population's own mean within-population distance divided by its cell count and therefore shifts candidates of different size or dispersion by different amounts. Every energy distance reported here is the U form except the PCA-latent population score of Fig. 2a, which is computed by a separate routine in the V form; the two forms are compared arm by arm in the code release. The population distances are computed on normalized expression rather than on control-referenced responses, whereas the mean-signature scores subtract the matched control mean. Fixed cell caps, penalties for undersized populations and the implementation settings for secondary scores are reported in Supplementary Methods.

Subpopulation-coverage score. The query is partitioned into K components ($k = 2$ throughout) and each candidate cell is assigned to the nearest query component. In the observed-population experiments the partition is the construction label of the two response modes and assignment is by Euclidean nearest centroid; in the partial-observation experiments, where no such label exists, the partition is k -means on the query cells with three initializations and a fixed seed. Writing d_k for the base distance between the query cells and the candidate cells of component k , the coverage distance and its ranking score are

$$D_\beta = \frac{1}{\beta} \left(\log \sum_{k=1}^K e^{\beta d_k} - \log K \right), \quad S_\beta = -D_\beta.$$

The base distance d_k is one of energy, MMD or sliced Wasserstein, so coverage is a family built on the same kernels rather than an independent metric. A candidate that populates no cell of a query component receives the penalty distance for that component, which is how the score charges a candidate for leaving part of the query uncovered.

As $\beta \rightarrow 0$, D_β averages subpopulation distances; as $\beta \rightarrow \infty$, it approaches the worst-matched component: 0.7125 and 1.300 for the four per-subpopulation distances drawn in Fig. 1h. For $K = 1$ it equals the global distance. Coverage-mean and coverage-worst are these two limits; numerical implementation details are given in Supplementary Methods.

Mean-signature baselines. For each response population $P_{d,c}$, the corresponding mean signature is $\bar{d}_{d,c} = \text{mean}(P_{d,c})$. Mean-signature retrieval ranks candidates using either cosine similarity or Euclidean distance between these vectors. The implemented CMap-style cosine baseline applies cosine similarity to the same control-referenced mean differential-expression vectors and is therefore numerically identical to mean-cosine retrieval (Spearman $\rho = 1.000$ over 54,180 query-candidate scores). The canonical L1000 weighted connectivity score (WTCS)¹² is included as a separate signature-based baseline. WTCS uses rank-based enrichment of the mean-collapsed expression response and is related to, but not algebraically identical to, the cosine implementation used here; it gives a different Hit@1 (0.4635).

To test whether the comparison depended on the original gene space, we also evaluated a PCA representation. One basis was fitted per retrieval query, on the query cells pooled with the cells of all 43 candidate populations of that query, the held-out correct candidate included, uniformly subsampled without replacement to 4,000 cells at a fixed seed and reduced to 30 components. The fit is therefore transductive but unsupervised: no candidate identity, ground-truth label or subpopulation label enters it, the correct candidate contributes on the same terms as the other 42, and the query cells are disjoint from the correct candidate's cells by construction. Control cells do not enter the fit, each control having already been collapsed to a mean vector that is projected through the fitted basis. The same comparison was then performed either between latent-space means or between the corresponding latent response populations.

Connection between mean and population retrieval. To formalize the relationship between the two representations, we write a response population as

$$X = \mu + \lambda\epsilon, \quad E[\epsilon] = 0,$$

where μ is the mean response, ϵ represents within-population variation and λ controls its magnitude. As $\lambda \rightarrow 0$, the response population collapses to a point mass at μ , and population comparison therefore loses all information beyond the mean. The mean *representation* is the endpoint of this variance-sensitivity construction. The score obtained at that endpoint is a distance between two mean vectors: for the energy distance between two point masses it is exactly $2\|\mu_Q - \mu_d\|_2$, which is the magnitude-aware mean score and not the cosine. Which score reads the mean is therefore a separate choice from how much of the population is retained, and the cosine is a further coarsening of the same representation that discards $\|\mu_Q - \mu_d\|_2$ altogether. Figure 2a measures what that coarsening costs.

The continuum refers to the amount of within-population response structure available to the scorer, not to an equivalence between otherwise different scoring functions: WTCS remains a distinct rank-based signature score, and energy, MMD and sliced Wasserstein retain their own distance definitions.

449 Response-matching and partial-observation evaluation

450 **Observed-population retrieval.** When complete candidate response populations were available, re-
451 trieval was evaluated over heterogeneous queries. Each query is a population built by combining two
452 response modes at mixing weight $\alpha \in \{0.5, 0.7, 0.9\}$, and the correct answer is the candidate that repro-
453 duces both modes at the same weight, built from the held-out half of the same cells so that no cell is
454 shared between a query and its answer. Two further candidates carry one mode each, and 40 distractor
455 populations complete a 43-candidate library.

456 The three tasks differ in what the two modes are. In the controlled SciPlex3 task they are the HDAC-
457 and JAK-associated response sources within a single cell line at 10 μ M, each pooling every compound
458 annotated to that class, so the two single-mode candidates are HDAC-only and JAK-only populations
459 and each distractor is an individual compound of another class. In the cross-line and Frangieh tasks
460 they are one perturbation observed in two biological contexts, two cell lines or Control and IFN γ , so the
461 two single-mode candidates are that perturbation observed in one context only and each distractor is a
462 different perturbation mixed across the same two contexts at the same weight; cross-line perturbations
463 pool all measured doses. The two single-mode candidates are each referenced to their own single-context
464 control, whereas the correct answer and every distractor are referenced to the α -mixed control.

465 Query perturbations in these two tasks are not arbitrary. Among perturbations that pass a split-half
466 reliability floor in both contexts and a cross-context response cosine below 0.9, the analysis uses the most
467 divergent: 12 drugs for each cell-line pair and 6 Frangieh knockouts. These two tasks therefore measure
468 retrieval where cross-context divergence is largest rather than on a representative perturbation panel.

469 Populations contain 400 cells in the controlled task and 200 in the other two, sampled without re-
470 placement where the source pool allows it and with replacement otherwise. Within one task setting
471 every query is drawn from a single fixed split of that setting’s cells and differs only in the perturbation
472 or mechanism sampled, in α and in the resampling seed, so queries within a setting are not independent
473 replicates. The controlled and cross-line tasks use ten resampling seeds and the Frangieh task five.

474 Hit@1 is reported as an unweighted macro-mean over the seven task–setting combinations, so a 30-
475 query combination counts as much as a 360-query one. Weighting instead by query count gives 0.884 for
476 energy distance and 0.421 for mean cosine, and PCA-space retrieval reaches 0.778 against 0.518 for the
477 corresponding mean-based score.

478 Hit@1 is the fraction of queries in which no candidate scores strictly higher than the both-mode
479 candidate; its rank is one plus the number of strictly higher-scoring candidates, and the same definition
480 underlies Hit@5, MRR and nDCG@10. That convention would resolve an exact tie in the both-mode can-
481 didate’s favour, but no tie arose: across the 1,260 queries and all nine scorers the both-mode candidate is
482 never tied with another candidate at any rank, so every Hit@1 reported here is identical under optimistic,
483 pessimistic and fractional tie handling. The partial-observation protocol below is scored differently: it
484 ranks by a label-blind stable sort and its Hit@1 is the best rank attained by any member of the query’s
485 mechanism or target set, not the rank of a single correct candidate. Ranking a single-mode candidate
486 first counts as a miss, including in the cross-line and Frangieh tasks, where that candidate carries the
487 correct perturbation observed in one of the two contexts. The criterion therefore scores the composition
488 of the retrieved population as well as the mechanism or perturbation it was drawn from. Comparisons
489 between population-level and mean-signature retrieval use the same queries and candidate libraries.
490 Aggregate Hit@1 values are reported as unweighted macro-means over the seven task–setting combina-
491 tions used in the observed-population analysis; query-weighted averages are reported as a sensitivity

analysis so that the largest task does not determine the primary summary.

Partial-observation protocol. We additionally evaluated retrieval when only part of each candidate response population was available to the retriever. This setting was designed to separate the information used to construct a ranking from the criterion subsequently used to evaluate that ranking. For each query, the population-level and mean-signature methods rank the same candidate library once; the resulting rankings are then retained unchanged while different evaluation criteria are applied.

Three splits construct the partial observation, all on SciPlex3, and in each the hidden material never enters the candidate library³⁷. In *leave-drug-out*, one drug is hidden, the query is that drug’s own treated population (up to 200 cells) and the library is every other drug in the context. In *leave-MoA-out*, an entire annotated mechanism class is removed from the library and the query is one drug drawn from the removed class, so recovery requires a same-target surrogate among the remaining drugs. In *partial-library*, a fraction of the drugs is treated as measured, the query is built from the pooled responses of the unmeasured drugs, and the retriever selects the next batch of candidates. In all three, query states are obtained by k -means with $k = 2$ on the query cells and the minority state is the smaller cluster; candidate populations are capped at 120–150 cells per drug.

The response-matching criterion is an energy-based welfare regret. Each candidate’s welfare is the aggregate over query states s of its negative energy distance to that state, $U(d) = \text{agg}_s(-\mathcal{E}(P_d, Q_s))$ with the aggregator taken as the minimum in the worst-case form, and the regret of the selected candidate $\hat{d}$ is

$$\text{Regret}(\hat{d}) = \max_{d \in \mathcal{L}} U(d) - U(\hat{d}) \geq 0,$$

where $\mathcal{L}$ is the observed library, so the reference is the best candidate actually available rather than an unobservable optimum. Minority-state coverage is reported on the same rankings as the cosine between the candidate’s mean response and the mean response of the query’s minority state, rescaled from $[-1, 1]$ to $[0, 1]$. Regret is zero when the top-ranked candidate is optimal under that objective and increases as the selected candidate becomes less favourable. We report regret reduction relative to mean-signature retrieval, such that a positive value indicates lower regret under this response-matching criterion. This evaluation is intentionally closely aligned with the population-level information used by the retriever and is therefore interpreted as evidence of response fidelity rather than as an independent functional outcome.

Task-proximal biological evaluation. The same fixed rankings were next evaluated using mechanism-of-action (MoA) recovery and minority-state coverage. MoA recovery was quantified by normalized discounted cumulative gain (nDCG)³⁸, which rewards candidates sharing the query perturbation’s annotated mechanism when they occur near the top of the ranked list. Mechanism recovery is scored on the 600 leave-drug-out queries, that is on the 765 partial-observation queries less the 120 leave-MoA-out and 45 partial-library queries, whose candidate libraries are built without a member of the query mechanism class and for which recovery is therefore undefined. Selection is by split type. Within the 600, 21 queries happen to hold no mechanism-class member either and score exactly zero under every method, so they enter the comparison as ties rather than as a difference; dropping them leaves the median difference at 0.000 and moves the mean from -0.024 to -0.025 and the decreased and increased shares from 40.8% and 35.7% to 42.3% and 37.0%.

Minority-state coverage asks whether the top-ranked candidate better preserves the response of the

less represented cellular state in the controlled heterogeneous-population construction. As with MoA recovery, this criterion is applied to rankings that have already been generated; the evaluator does not alter the query, the retrieval score or the candidate ordering. Differences between population-level and mean-signature retrieval are therefore attributable to how the same rankings score under the biological criterion rather than to refitting or re-ranking.

Cross-setting projection. To examine whether the minority-state result extends beyond the controlled partial-observation setting, we evaluated the top-ranked population-level and mean-signature candidates across 239 retrieval task instances spanning natural and constructed settings. These comprised 90 constructed cross-cell-line SciPlex3 tasks, 90 within-cell-line SciPlex3 tasks, 23 Frangieh tasks, 12 CD34+ tasks and 24 tasks using predicted candidate populations. For each task we report the difference in minority-state coverage between the population-level and mean-signature selections.

The 239 tasks are treated as a distribution of task-level differences rather than as independent biological replicates of a single experiment. In particular, the constructed SciPlex3 tasks share data source, preprocessing and candidate structure, and the three SciPlex3 cell lines constitute leave-one-cell-line-out variation within one experiment rather than leave-one-dataset-out generalization. Thresholded counts of “distributionally dominant” tasks are therefore not used as a primary result; the continuous difference in coverage is reported instead.

External functional evaluation using GDSC2

To test whether transcriptomic retrieval rankings capture information that extends beyond the single-cell expression measurements used to construct them, we evaluated the SciPlex3 rankings against drug-response profiles from the Genomics of Drug Sensitivity in Cancer^{28,29}, using the GDSC2 screening set (release 8.5; Data availability). The external criterion was defined as similarity between the functional response profiles of two drugs rather than as absolute drug potency, because the retrieval task asks which candidate produces a response most similar to the query rather than which candidate is most cytotoxic.

SciPlex3 compounds were matched to GDSC2 by normalized drug name, with substring matching used only as a fallback. After removal of one duplicate mapping and restriction to the GDSC2 screening platform, 35 compounds remained. GDSC1 and GDSC2 measurements were not pooled because they derive from different assay and dose-range settings. For each matched compound, its functional profile was the vector of fitted dose-response AUC values across GDSC2 cell lines. The three cell lines used for the transcriptomic retrieval experiments, A549, K562 and MCF7, were excluded from these profiles, leaving 966 cell lines, so the functional evaluation was performed in biological contexts not used to construct the RNA rankings.

Cell lines differ systematically in their overall sensitivity to drug treatment, which can induce positive correlations between otherwise unrelated compound profiles. We therefore removed this shared cell-line effect before calculating drug-drug similarity. For cell line l and compound d , the centred response was

$$\widetilde{\text{AUC}}_{d,l} = \text{AUC}_{d,l} - \frac{1}{|\mathcal{D}_l|} \sum_{d' \in \mathcal{D}_l} \text{AUC}_{d',l},$$

where $\mathcal{D}_l$ contains all GDSC2 compounds measured in cell line l , rather than only the 35 compounds overlapping SciPlex3. Functional similarity between query drug q and candidate drug c was then defined as

$$F(q, c) = \rho_s \left(\widetilde{\text{AUC}}_{q,\bullet}, \widetilde{\text{AUC}}_{c,\bullet} \right),$$

where ρ_S denotes Spearman correlation across cell lines shared by the two compounds. Drug pairs measured in fewer than 100 common cell lines were excluded. This yielded 591 usable drug pairs, with a median of 909 shared cell lines per pair. The corresponding uncentred functional similarity was retained as a sensitivity analysis.

For each SciPlex3 cell line, retrieval was evaluated in a leave-one-drug-out design. The 10 μ M response population of the held-out compound served as the query, and the remaining matched compounds formed the candidate library, giving 34–35 candidates per query. Each RNA-based retrieval rule produced a candidate ranking without access to GDSC2 data. Agreement between that ranking and the external functional ordering $F(q, c)$ was quantified by Spearman correlation across candidates.

The primary RNA-derived rankings were energy-distance population retrieval and control-referenced mean-cosine retrieval. To test whether apparent functional agreement required cell-to-cell distributional information, we additionally used a response-magnitude matching score,

$$M(q, c) = - \left| \|\bar{d}_q\| - \|\bar{d}_c\| \right|,$$

which prefers candidates whose overall transcriptional response magnitude is similar to that of the query but uses no within-population structure. A query-independent magnitude score $\|\bar{d}_c\|$ was retained as a diagnostic control. An uncorrected mean-cosine score based on treated means without matched-control subtraction and a potency-matching diagnostic were analysed as additional sensitivity checks rather than as primary retrieval baselines.

To estimate the component of the energy-based association not explained by response magnitude, we computed a partial rank correlation with the magnitude-matching channel removed. Retrieval scores and functional similarities were rank-transformed, each was residualized with respect to $M(q, c)$, and the correlation between the two residuals was reported as the magnitude-adjusted association.

Per-query Spearman correlations are summarized by their median within each SciPlex3 cell line. Queries from the same cell line share the same candidate library and are therefore not treated as independent biological replicates. The three cell lines provide the higher-level analysis unit; query-level paired tests are reported only as descriptive analyses where indicated.

The comparison between energy retrieval and the response-magnitude scalar is pre-specified as a pooled test across the three SciPlex3 cell lines; per-line ratios are given in Supplementary Methods.

Protein-based evaluator-shape experiment

Frangieh Perturb-CITE-seq³⁰ provides matched RNA and surface-protein measurements from the same cells, allowing the statistical form of an external evaluator to be changed while holding its biological measurement fixed. RNA and protein matrices were joined by cell barcode before analysis; row position was not used because the released modality-specific matrices differ in cell number and ordering. Of the 24 measured antibodies, four isotype controls were removed after normalization, leaving 20 surface-protein features. Antibody-derived tag counts³⁹ were centred-log-ratio normalized within each cell, $y_{ij} = \ln(1 + x_{ij}) - m^{-1} \sum_k \ln(1 + x_{ik})$ for cell i and antibody j , with a pseudocount of one, a natural logarithm and the per-cell mean taken over all $m = 24$ measured antibodies including zero counts and the four isotype controls; normalization therefore precedes isotype removal. This is the centred log-ratio of the $(1 + x)$ vector and differs from implementations that exclude zero entries from the geometric mean.

All analyses were performed separately within the Control, IFN γ and Co-culture conditions to avoid

610 mixing immune context with perturbation identity. Knockouts represented by at least 60 barcode-
 611 matched cells were retained, and populations were subsampled to at most 250 cells using a fixed random
 612 seed.

613 **Protein-based evaluators.** For knockout k in immune condition c , let $\bar{d}_{k,c}^{\text{prot}}$ denote its mean control-
 614 referenced protein response and let $D_{k,c}^{\text{prot}}$ denote the corresponding set of per-cell protein responses. The
 615 mean-based evaluator was defined as

$$E_{\text{mean}}(q, d) = \cos \left(\bar{d}_{q,c}^{\text{prot}}, \bar{d}_{d,c}^{\text{prot}} \right).$$

616 The population-based evaluator was defined from the same cells and protein measurements as

$$E_{\text{pop}}(q, d) = -\mathcal{E} \left(D_{q,c}^{\text{prot}}, D_{d,c}^{\text{prot}} \right),$$

617 where $\mathcal{E}$ is the energy distance. The negative sign converts distance to similarity so that larger values
 618 indicate a closer match under both evaluators.

619 For both definitions, protein responses were referenced to the matched control population from the
 620 same immune condition. Thus,

$$\bar{d}_{k,c}^{\text{prot}} = \text{mean} \left(X_{k,c}^{\text{prot}} \right) - \mu_{0,c}^{\text{prot}},$$

621 and

$$D_{k,c}^{\text{prot}} = \left\{ x_{i,k,c}^{\text{prot}} - \mu_{0,c}^{\text{prot}} \right\}_{i=1}^{n_{k,c}}.$$

622 Using the same matched control for the two evaluator forms ensures that their difference arises from
 623 whether protein responses are summarized by their mean or retained as a cellular population, rather
 624 than from a change in biological reference.

625 **Fixed RNA-derived rankings.** The RNA-derived rankings were computed once and then evaluated
 626 against both protein criteria. The primary rankings were energy-distance population retrieval and
 627 control-referenced mean-cosine retrieval. The query populations, candidate libraries and RNA-derived
 628 candidate scores were identical under the two protein evaluators, and no protein measurement was
 629 available to any retrieval method. The experiment therefore changes only the statistical representation
 630 used to judge the same RNA-derived ranking.

631 **Response-magnitude control.** Because a population distance can also encode differences in overall re-
 632 sponse magnitude, we included a query-dependent scalar control that uses no cell-to-cell distributional
 633 information. For query q and candidate d , magnitude matching was defined as

$$M(q, d) = - \left| \left| \bar{d}_{q,c}^{\text{RNA}} \right| - \left| \bar{d}_{d,c}^{\text{RNA}} \right| \right|.$$

634 This ranking favours candidates with a transcriptional response magnitude similar to that of the query
 635 while discarding all within-population structure. The same magnitude ranking was evaluated against
 636 both protein criteria.

637 For each query, agreement between an RNA-derived ranking and a protein-based evaluator was
 638 quantified by Spearman correlation across that query's candidate knockouts. The statistic is therefore a

rank correlation over candidates, computed once per query, and never a correlation over pooled query-candidate pairs. The headline value for each ranking and each evaluator form is the median of those per-query correlations over all 659 queries. The same medians were also taken within each of the three immune conditions, which are reported separately and treated as the higher-level analysis unit; queries within a condition share candidate structure and were not interpreted as independent biological replicates.

Natural heterogeneity and large-scale validation

We next examined the information requirements for population-level retrieval in datasets in which cellular heterogeneity was not created by mixing experimentally distinct source populations. Patient-derived glioblastoma was used to test the requirements in naturally heterogeneous tissue, and Tahoe-100M was used to assess the same quantities at substantially larger scale. The glioblastoma cohort was not used as a retrieval benchmark; Tahoe-100M additionally served as the substrate for the source-to-target intervention-identification benchmark described below.

Patient-derived glioblastoma. The ZhaoSims2021 dataset¹⁶ contains acute-slice cultures and surgical biopsies from ten glioblastoma patients exposed to six compounds: Ana-12, Ispenisib, RO4929097, Tazemetostat, etoposide and panobinostat. Expression counts were normalized to 10^4 per cell, log1p-transformed and reduced to 2,000 highly variable genes. Because only one patient received more than two compounds, the cohort does not support a sufficiently replicated leave-one-drug-out retrieval benchmark. We therefore use it only to measure state-specific response, recoverability and the extent to which population-level similarity differs from compartment-specific response ordering.

Cellular-compartment assignment. Cells were assigned to malignant and myeloid compartments using canonical marker scores; ambiguous cells were excluded. Drug responses were defined relative to patient- and compartment-matched controls. Assignment criteria, retained-cell counts and threshold-sensitivity analyses are reported in Supplementary Note 3.

Differential response in natural tissue. State-specific perturbation responses were measured between malignant and myeloid compartments. For patient p , drug d and compartment k , the response vector was

$$r_{p,d,k} = \text{mean} \left(X_{p,d,k}^{\text{treated}} \right) - \text{mean} \left(X_{p,k}^{\text{control}} \right).$$

Only patient-drug combinations with at least 50 treated cells in both compartments were included. Controls were matched within both patient and compartment. Differential response was summarized as

$$R_{p,d} = \cos \left(r_{p,d,\text{malignant}}, r_{p,d,\text{myeloid}} \right).$$

This construction removes pre-existing differences between patients and cellular compartments before comparing the directions of their drug-induced responses.

Recoverability in natural tissue. Response-identity recoverability was assessed from within-patient drug-pair splits by comparing unsupervised partitions with a supervised ceiling. Patient-level bootstrap and method-pair sensitivity analyses are reported in Supplementary Note 3.

673 **Proxy for downstream decision relevance.** The glioblastoma cohort does not contain enough drug
 674 coverage to test whether population-level retrieval directly changes the preferred candidate. We there-
 675 fore quantified a more limited property: whether drug similarity at the whole-population level preserves
 676 the ordering observed within a specific cellular compartment.

677 For each pair of drugs measured in the same patient, we first computed the cosine similarity be-
 678 tween their population-mean response signatures and the cosine similarity between their malignant-
 679 compartment response vectors. Agreement between these two similarity structures was summarized by
 680 Spearman correlation across drug pairs.

681 That whole-population comparison gives a Spearman correlation of +0.878, but the population-mean
 682 signature is the equal-weight average of the two compartment responses, so the malignant compartment
 683 contributes half of it and part of the agreement is arithmetic. We therefore additionally used a disjoint
 684 comparison. For drug pair (a, b) , myeloid-response similarity was defined as

$$S_{a,b}^{\text{myeloid}} = \cos(r_{a,\text{myeloid}}, r_{b,\text{myeloid}}),$$

685 and malignant-response similarity as

$$S_{a,b}^{\text{malignant}} = \cos(r_{a,\text{malignant}}, r_{b,\text{malignant}}).$$

686 We then computed

$$\rho_{\text{disjoint}} = \rho_S(S_{a,b}^{\text{myeloid}}, S_{a,b}^{\text{malignant}}),$$

687 using only drug pairs available within the same patient. The two sides of this comparison contain no
 688 shared cells and each compartment is referenced to its own matched control. This statistic is interpreted
 689 as a proxy for the amount of candidate-ordering headroom available to population structure, not as a
 690 direct measure of retrieval performance. No pass-fail threshold was applied.

691 **Large-scale validation in Tahoe-100M.** Tahoe-100M plate 3 was used to test the same requirements at
 692 larger scale. We evaluated two independently defined states, released cell-cycle labels and a control-
 693 derived k -means partition, and quantified state-specific response alignment, response-identity recover-
 694 ability and agreement of drug ordering across disjoint states. Dataset filtering and all state-definition
 695 details are reported in Supplementary Note 3.

696 For each eligible drug-cell-line condition, the response of state k was defined relative to its own
 697 state-matched controls,

$$r_{d,c,k} = \text{mean}(X_{d,c,k}^{\text{treated}}) - \text{mean}(X_{c,k}^{\text{control}}).$$

698 Differential response was measured as the cosine between the two state-specific response vectors,

$$R_{d,c} = \cos(r_{d,c,0}, r_{d,c,1}).$$

699 A condition was included only when at least 50 cells were available in each required treated and control
 700 state.

701 Finally, a disjoint state-ordering statistic was used as a large-scale proxy for downstream decision
 702 relevance. Within each cell line, cosine similarity was computed for every eligible pair of drug responses

separately in the two non-overlapping cellular states. Spearman correlation between the two resulting drug-drug similarity matrices was then calculated,

$$\rho_c = \rho_S \left(S_{a,b}^{(0)}, S_{a,b}^{(1)} \right).$$

High correlation indicates that the two cellular states preserve similar relative ordering among drug responses, whereas lower correlation leaves greater scope for state-specific information to change candidate preference. As in the glioblastoma analysis, this quantity does not directly test whether population-level retrieval improves a drug ranking.

These are intrinsic-state diagnostics and are reported separately from the intervention-identification benchmark below, which uses the same plate. Detailed recoverability methods are in Supplementary Note 3.

Source-to-target intervention identification on Tahoe-100M

The retrieval experiments above rank candidates whose responses in the query context are observed. An intervention task instead asks which compound will carry a starting population into a desired one, so the candidate’s effect in that context is unknown. The scope of that question here is narrower than intervention design, and the naming follows the scope. The desired population is a real treated population generated by one compound of the library, and the outcome is the rank that compound receives, so the task measures whether population information helps identify the intervention that produced an observed response. It does not measure therapeutic utility, and no target state is specified independently of a compound that already reaches it. We defined this task on Tahoe-100M plate 3.

Task definition and eligibility. A context is one cell line. For a held-out context c and a query drug q , the retriever receives the vehicle (DMSO) cells of c and an observed treated population $X_{q,c}$, and ranks all 92 candidate compounds; it never sees q . A condition entered the grid if it had at least 100 treated and 200 vehicle cells, which admitted 3,992 eligible queries over 44 usable cell lines from the 4,600 conditions present on the plate. Each query was evaluated at five fixed seeds (13, 29, 47, 71, 101), giving 19,960 query-seed pairs. Source and target populations were drawn without replacement at up to 200 cells each, all available cells being taken where a condition held fewer, and responses were referred to the source mean, $P = \{x_i - \bar{x}_{\text{source}}\}$.

Oracle arm. In the oracle arm every candidate’s response in the held-out context is taken from the observed data, which removes the forward model from the loop and bounds what any predictor could deliver. Cell disjointness is enforced through the library preparation: the target population is drawn from one sublibrary half and every candidate bank from the other, at up to 100 cells per candidate, all cells in that half being taken where a candidate has fewer, without replacement, so the generating drug is drawn from the same half as its 91 competitors and is not penalized by a technical shift. This arm is an information ceiling and not a deployment result.

Prediction arm. In the prediction arm no treated cell of the held-out context enters any fit or any ranker. Candidate responses are produced by a forward predictor fitted on the other 43 cell lines. Six predictors were evaluated. Three are additive: they emit one signature per drug and apply it identically to every source cell, so the candidate population is the vehicle population rigidly translated. These are an average effect δ_d pooled over training contexts; a nearest-context transfer taking δ_d from the training context

whose vehicle mean is closest in cosine; and a linear-latent variant truncating the average effect to the principal components carrying 80% of training variance. Three are not additive. The *state-conditioned* predictor estimates the same average effect separately for each of the three cell-cycle states released with the plate, G1, S and G2M, and applies it to each source cell according to its own state, which costs one extra grouping and no new hyperparameter. Each per-state effect is averaged over the training contexts in which that state holds at least 20 cells; groups below the floor are dropped from the average rather than estimated, and 359 of the 12,275 cell-line, drug and state groups fall below it, 294 of them in S phase. A drug left with no usable estimate for a state would keep its pooled effect there and be additive for that state alone; this did not occur for any of the $44 \times 92 \times 3$ signature slots, so the predictor is non-additive over its whole output. The *Gaussian optimal-transport* map is a Bures map fitted in a 50-component training principal-component basis, from the pooled within-context vehicle covariance to the drug’s pooled within-condition treated covariance, with the mean shifted by the average latent effect and applied to the held-out vehicle cells as a latent edit so that the part of each cell outside the subspace is carried through; every cloud is centred on its own condition mean before pooling, so a covariance here is within-condition cell scatter. *Nearest context, cells* transfers the neighbouring context’s actual treated cells for the drug, shifted by the difference of the two vehicle means; it is the cheapest predictor in this set that can carry a non-additive response shape, and it is included because a set of purely additive predictors would answer the question by construction rather than by measurement. Each principal-component basis is fitted on up to 50,000 cells sampled from the training contexts, and each context–drug fit uses up to 200 cells. The additive three generate no state-specific differential response, which is an algebraic property of the model class and not a training failure.

Scoring and evaluation. Candidates were scored by the three rules of Fig. 1a: a direction-only mean cosine, a magnitude-aware mean L_2 , and the population route, primarily the unbiased U-statistic energy distance, with MMD and sliced Wasserstein as secondary variants. Mean reciprocal rank of the true intervention is the primary outcome; Hit@1, Hit@5 and Hit@10 are secondary. Intervals are 2,000-fold bootstrap intervals resampled over cell-line clusters, which are the independent units of this design. A decision is counted as corrected when the population route places the true drug first and the reference does not, and as broken in the reverse case.

Headroom, response detectability and the conditional bottleneck analysis

Headroom and its null. Headroom is defined per query as $H = 1 - \text{RR}_{\text{ref}}$, the room the direction-only mean route leaves for any other scorer to improve on, and the gain is $G = \text{RR}_{\text{pop}} - \text{RR}_{\text{ref}}$ on the same query. A Spearman correlation between H and G is partly arithmetic, because the two share RR_{ref} with opposite signs and reciprocal rank is capped at 1. The null quantifies exactly that. Within each cell line, the population route’s reciprocal ranks are randomly permuted across queries, which destroys any pairing between the two scorers while preserving both marginal distributions, the $\text{RR} \leq 1$ ceiling and the cell-line clustering; G is then recomputed against the unpermuted reference and the correlation retaken. Two hundred permutations at a fixed seed give the interval reported in Fig. 5e.

Response detectability. For each held-out cell line c and query drug q , equal numbers of cells treated with q and vehicle cells of c form a balanced two-class problem, at 300 per arm or at the size of the smaller pool where that is below 300, sampled at a fixed seed, so chance is 0.5 by construction. The supervised ceiling A_{sup} is the out-of-fold accuracy of a five-fold cross-validated logistic probe ($C = 1$) on those cells.

782 The unsupervised value A_{unsup} is k -means with $k = 2$ on the same cells, scored by best-permutation
 783 accuracy, so cluster-to-label matching is free and the unsupervised arm is not penalized for arbitrary
 784 label order. The labels are external, being which well a cell came from, rather than an expression-derived
 785 cell state. The supervised arm is a ceiling on the information present in the representation.

786 **Conditional bottleneck regression.** Regressing the gain G directly on candidate explanators inherits
 787 the same arithmetic as the headroom correlation. The specification reported in Fig. 5h therefore takes
 788 the population route’s own reciprocal rank as the outcome and includes the mean route’s reciprocal
 789 rank as a covariate, so that a coefficient answers a conditional question: given how well the mean route
 790 did on this query, does this explainer predict how well the population route did. The covariate is
 791 a control and is not plotted. The explanators are the interaction share, its cross-fitted reproducibility
 792 $R_{\text{int}} = \cos(I_A, I_B)$, response detectability A_{sup} , and, in a separate model reported for contrast, the naive
 793 response cosine $1 - \cos(r_1, r_2)$. Estimation is ordinary least squares with cluster-robust standard errors
 794 at the cell line, which is the independent unit of the design, and all predictors are standardised to zero
 795 mean and unit variance so that coefficients are comparable across explanators on one axis. Queries
 796 missing any explainer are dropped listwise, which is why the reported n of 3,618 is below the 3,992 of
 797 the retrieval arm.

798 **Cross-fitted drug-by-state interaction**

799 Testing whether a drug’s effect depends on cell state requires a statistic that measures interaction rather
 800 than effect size. Writing r_s for the state- s response, the difference between mean treated and mean vehicle
 801 cells of that state, with states 1 and 2 the G1 and G2M cell-cycle calls released with the plate, the quantity
 802 of interest is the interaction vector $I = r_1 - r_2$. Its squared norm cannot be used directly: with finite arms,
 803 sampling noise alone makes $\|I\|$ large, and the same mechanism inflates $1 - \cos(r_1, r_2)$ for the weakest
 804 drugs.

805 Each of the four arms is therefore split into disjoint halves A and B , the interaction is estimated once
 806 in each, and the two estimates are multiplied:

$$S_{\text{int}} = \frac{\langle I_A, I_B \rangle}{p}, \quad S_{\text{main}} = \frac{\langle M_A, M_B \rangle}{p}, \quad M = \frac{1}{2}(r_1 + r_2),$$

807 with p the number of genes. Under the additive null $E[I] = 0$ and the two halves are independent,
 808 so $E[S_{\text{int}}] = 0$ exactly: noise makes I_A and I_B point in unrelated directions and their inner product
 809 averages to zero, whereas a real state-dependent effect survives because both halves estimate the same
 810 non-zero I . The vehicle arms are split as well; sharing one vehicle mean between halves would place the
 811 same control noise in I_A and I_B and correlate them under the null. Splits are library-preparation-disjoint
 812 wherever a condition spans at least two sublibraries, and a plain random half of the cells otherwise.
 813 Arms were capped at 400 cells and required at least 50, and ten A/B splits were averaged per query.

814 Two derived quantities are used. The *interaction share*, $S_{\text{int}}/(S_{\text{int}} + S_{\text{main}})$, is scale-free and has a
 815 cross-fitted numerator and denominator, so neither term is noise-inflated; dividing by a plain response
 816 norm would reintroduce exactly the single-sample denominator that sank the cosine statistic. $R_{\text{int}} =$
 817 $\cos(I_A, I_B)$ asks whether the interaction reproduces across halves. As a positive control, synthetic condi-
 818 tions were built from the plate’s own cells by taking the G1 cells from one drug and the G2M cells from
 819 another, so interaction is present by construction in material with the same states and arm sizes as the
 820 queries.

821 The cosine statistic $1 - \cos(r_1, r_2)$ is retained only as a descriptive sensitivity analysis and as the his-
822 torical comparator in Fig. 5g. It is not used as a differential-response measure anywhere that a conclu-
823 sion depends on it. Where earlier analyses in this study used the induced-response cosine to characterize
824 predictors or tissue, that use is descriptive alignment between two response directions and is labelled as
825 such.

826 Statistical analysis

827 Statistical units follow the experimental design rather than the number of cells or query–candidate com-
828 parisons: cells estimate a perturbation-response population but are not independent replicates of re-
829 trieval performance⁴⁰.

830 Paired comparisons between retrieval methods on the same queries used two-sided Wilcoxon signed-
831 rank tests; independent groups used Mann–Whitney U tests; monotonic associations used Spearman
832 rank correlation. Across pre-specified strata, P values were adjusted by the Benjamini–Hochberg pro-
833 cedure⁴¹. Effect sizes are reported with every test; achieved power and the sample size required for
834 80% power are computed for the one stratum where they bear on the conclusion and are reported in
835 Supplementary Note 1. Mechanism-of-action analyses were restricted to the 600 queries whose mecha-
836 nism class remained represented in the candidate library; undefined queries were excluded rather than
837 encoded as zero.

838 For the GDSC2 functional analysis, Spearman correlations were computed per query across its candi-
839 date drugs and summarized within each SciPlex3 cell line, which is the analysis unit; query-level paired
840 tests are reported as nominal. The protein evaluator-shape experiment was summarized within the Con-
841 trol, IFN γ and Co-culture conditions, its three higher-level units. For HIR-Bench, cross-validation was
842 grouped by the label-determining (α , conflict) pair, giving 28 effective parameter cells, and performance
843 was summarized by pooled out-of-fold AUC. For the patient-derived glioblastoma analysis, uncertainty
844 in the recoverability gap was estimated by cluster bootstrap at the patient level with 5,000 resamples, dif-
845 ferences being taken within each drug-pair split before aggregation; four patients contributed evaluable
846 splits. Intervention-retrieval intervals are 2,000-fold bootstrap resamples over cell-line clusters. Tahoe-
847 100M summaries are taken across eligible drug–cell-line conditions or within-cell-line drug pairs accord-
848 ing to the statistic, and both cellular-state definitions, released cell-cycle state and the control-derived
849 k -means partition, were analysed in parallel.

850 Analyses were implemented in Python. Single-cell data handling used Scanpy⁴²; the k -means,
851 Gaussian-mixture, principal-component and logistic-regression procedures named above are the
852 scikit-learn implementations⁴³.

853 Data availability

854 All datasets analysed in this study are publicly available.

855 SciPlex3, Frangieh Perturb-CITE-seq and ZhaoSims2021 were obtained in harmonized form
856 from scPerturb¹³, Zenodo record 13350497 ([doi:10.5281/zenodo.13350497](https://doi.org/10.5281/zenodo.13350497), release 1.4). The exact
857 harmonized files used in the analyses were SrivatsanTrapnell2020_sciplex3.h5ad for SciPlex3,
858 FrangiehIzar2021_RNA.h5ad and FrangiehIzar2021_protein.h5ad for Frangieh Perturb-CITE-seq, and
859 ZhaoSims2021.h5ad for the patient-derived glioblastoma analysis. The primary SciPlex3 data are avail-
860 able from the Gene Expression Omnibus under accession GSE139944¹⁴; the original Frangieh and
861 ZhaoSims2021 datasets are described in the corresponding publications^{16,30}. The CD34+ haematopoietic
862 progenitor dataset is available from the Gene Expression Omnibus under accession GSE306429³².

863 Tahoe-100M¹⁵ was released by Vevo Therapeutics and the Arc Institute under CC0 and is distributed
864 through the Arc Virtual Cell Atlas and Hugging Face (tahoebio/Tahoe-100M). This study uses plate 3
865 only, together with the released obs_metadata and sample_metadata tables. Analyses use the released
866 preprocessed expression matrix comprising 4,158,278 cells and 2,304 highly variable genes in log1p-
867 transformed space.

868 The external functional evaluation used GDSC2 fitted dose-response data from the Genomics of
869 Drug Sensitivity in Cancer, release 8.5 (GDSC2_fitted_dose_response_270ct23.xlsx), available from
870 <https://www.cancerrxgene.org>. The compound-matching table used to align the SciPlex3 and GDSC2
871 drug panels (drug_match_table.csv) is provided with the analysis code.

872 Processed analysis inputs and the scripts used to construct them from the source datasets are pro-
873 vided in the accompanying repository. Source-data tables underlying each figure panel are provided
874 with the corresponding analysis outputs.

875 Code availability

876 All code used for data processing, retrieval, evaluation and statistical analysis is available at <https://github.com/Boom5426/PopRetrieve>.
877

878 Acknowledgements

879 **[[Acknowledgements. Funding sources and grant numbers.]]**

880 Author contributions

881 **[[Author contributions, CRediT taxonomy: conceptualization, methodology, software, formal analy-
882 sis, investigation, data curation, writing (original draft), writing (review and editing), visualization,
883 supervision, funding acquisition.]]**

884 Competing interests

885 **[[Competing interests statement. If none: "The authors declare no competing interests."]]**

886 Additional information

887 Supplementary Information is available for this paper: Supplementary Methods, Supplementary Notes
888 1–4 and Supplementary Tables 1–4.

889 Correspondence and requests for materials should be addressed to **[[Name]]**.

890 References

- 891 [1] Charlotte Bunne, Yusuf Roohani, Yanay Rosen, et al. How to build the virtual cell with artificial
892 intelligence: priorities and opportunities. *Cell*, 187(25):7045–7063, 2024. doi: 10.1016/j.cell.2024.11.
893 015.
- 894 [2] Yusuf H. Roohani et al. Virtual Cell Challenge: toward a Turing test for the virtual cell. *Cell*, 188
895 (13):3370–3374, 2025. doi: 10.1016/j.cell.2025.06.008.
- 896 [3] Haotian Cui, Chloe Wang, Hassaan Maan, Kuan Pang, Fengning Luo, Nan Duan, and Bo Wang.
897 scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature*
898 *Methods*, 21(8):1470–1480, 2024. doi: 10.1038/s41592-024-02201-0.
- 899 [4] Minsheng Hao et al. Large-scale foundation model on single-cell transcriptomics. *Nature Methods*,
900 21(8):1481–1491, 2024. doi: 10.1038/s41592-024-02305-7.

- [5] Christina V. Theodoris et al. Transfer learning enables predictions in network biology. *Nature*, 618 (7965):616–624, 2023. doi: 10.1038/s41586-023-06139-9.
- [6] Mohammad Lotfollahi, F. Alexander Wolf, and Fabian J. Theis. scGen predicts single-cell perturbation responses. *Nature Methods*, 16(8):715–721, 2019. doi: 10.1038/s41592-019-0494-8.
- [7] Mohammad Lotfollahi et al. Predicting cellular responses to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19(6):e11517, 2023. doi: 10.15252/msb.202211517.
- [8] Leon Hetzel, Simon Böhm, Niki Kilbertus, Stephan Günnemann, Mohammad Lotfollahi, and Fabian J. Theis. Predicting cellular responses to novel drug perturbations at a single-cell resolution. In *Advances in Neural Information Processing Systems*, volume 35, 2022. NeurIPS 2022 proceedings register no DOI.
- [9] Yusuf Roohani, Kexin Huang, and Jure Leskovec. Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology*, 42(6):927–935, 2024. doi: 10.1038/s41587-023-01905-6.
- [10] Guadalupe Gonzalez, Xiang Lin, Isuru Herath, Kirill Veselkov, Michael Bronstein, and Marinka Zitnik. Combinatorial prediction of therapeutic perturbations using causally inspired neural networks. *Nature Biomedical Engineering*, 10(5):1008–1025, 2026. doi: 10.1038/s41551-025-01481-x.
- [11] Justin Lamb et al. The Connectivity Map: using gene-expression signatures to connect small molecules, genes, and disease. *Science*, 313(5795):1929–1935, 2006. doi: 10.1126/science.1132939.
- [12] Aravind Subramanian et al. A next generation Connectivity Map: L1000 platform and the first 1,000,000 profiles. *Cell*, 171(6):1437–1452.e17, 2017. doi: 10.1016/j.cell.2017.10.049.
- [13] Stefan Peidli et al. scPerturb: harmonized single-cell perturbation data. *Nature Methods*, 21(3): 531–540, 2024. doi: 10.1038/s41592-023-02144-y.
- [14] Sanjay R. Srivatsan et al. Massively multiplex chemical transcriptomics at single-cell resolution. *Science*, 367(6473):45–51, 2020. doi: 10.1126/science.aax6234.
- [15] Jesse Zhang et al. Tahoe-100M: a giga-scale single-cell perturbation atlas for context-dependent gene function and cellular modeling. *bioRxiv*, 2025. doi: 10.1101/2025.02.20.639398. Preprint, version 3 (2025-05-10).
- [16] Wenting Zhao et al. Deconvolution of cell type-specific drug responses in human tumor tissue with single-cell RNA-seq. *Genome Medicine*, 13(1):82, 2021. doi: 10.1186/s13073-021-00894-y.
- [17] Sydney M. Shaffer et al. Rare cell variability and drug-induced reprogramming as a mode of cancer drug resistance. *Nature*, 546(7658):431–435, 2017. doi: 10.1038/nature22794.
- [18] Yuge Ji, Tessa D. Green, Stefan Peidli, Mojtaba Bahrami, Meiqi Liu, Luke Zappia, Karin Hrovatin, Chris Sander, and Fabian J. Theis. Optimal distance metrics for single-cell RNA-seq populations. *bioRxiv*, 2023. doi: 10.1101/2023.12.26.572833. preprint, bioRxiv.
- [19] Andrea Rubbi, Andrea Giuseppe Di Francesco, Mohammad Lotfollahi, and Pietro Liò. A standardized framework for evaluating gene expression generative models. *arXiv*, 2026. doi: 10.48550/arXiv.2603.11244. preprint arXiv:2603.11244.

- [20] Charlotte Bunne et al. Learning single-cell perturbation responses using neural optimal transport. *Nature Methods*, 20(11):1759–1768, 2023. doi: 10.1038/s41592-023-01969-x.
- [21] Dominik Klein et al. CellFlow enables generative single-cell phenotype modeling with flow matching. *bioRxiv*, 2025. doi: 10.1101/2025.04.11.648220. preprint, bioRxiv.
- [22] Changxi Chi, Yufei Huang, Jun Xia, Jiangbin Zheng, Yunfan Liu, Zelin Zang, and Stan Z. Li. Departures: distributional transport for single-cell perturbation prediction with neural Schrödinger bridges. *Proceedings of the AAAI Conference on Artificial Intelligence*, 40(25):20544–20552, 2026. doi: 10.1609/aaai.v40i25.39190. also available as arXiv:2511.13124.
- [23] Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep-learning-based gene perturbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 22(8):1657–1661, 2025. doi: 10.1038/s41592-025-02772-6.
- [24] Kasia Z. Kedzierska, Lorin Crawford, Ava P. Amini, and Alex X. Lu. Zero-shot evaluation reveals limitations of single-cell foundation models. *Genome Biology*, 26(1):101, 2025. doi: 10.1186/s13059-025-03574-x.
- [25] Zhiting Wei et al. Benchmarking algorithms for generalizable single-cell perturbation response prediction. *Nature Methods*, 23(2):451–464, 2026. doi: 10.1038/s41592-025-02980-0.
- [26] Mahshid Heidari, Mina Karimpour, Sumana Srivatsa, and Hesam Montazeri. Evaluating single-cell perturbation response models is far from straightforward. *bioRxiv*, 2026. doi: 10.64898/2026.02.14.705879. preprint, bioRxiv.
- [27] Rachel L. Thomas and David Uminsky. Reliance on metrics is a fundamental challenge for AI. *Patterns*, 3(5):100476, 2022. doi: 10.1016/j.patter.2022.100476.
- [28] Wanjuan Yang et al. Genomics of Drug Sensitivity in Cancer (GDSC): a resource for therapeutic biomarker discovery in cancer cells. *Nucleic Acids Research*, 41(D1):D955–D961, 2013. doi: 10.1093/nar/gks1111.
- [29] Francesco Iorio et al. A landscape of pharmacogenomic interactions in cancer. *Cell*, 166(3):740–754, 2016. doi: 10.1016/j.cell.2016.06.017.
- [30] Chris J. Frangieh et al. Multimodal pooled Perturb-CITE-seq screens in patient models define mechanisms of cancer immune evasion. *Nature Genetics*, 53(3):332–341, 2021. doi: 10.1038/s41588-021-00779-1.
- [31] Dhruv Agarwal and Riya Bisht. The metric picks the winner: evaluation choice flips model rankings for drug-response prediction in unseen chemistry. *arXiv*, 2026. doi: 10.48550/arXiv.2606.12639. preprint arXiv:2606.12639.
- [32] Benjamin DeMeo et al. Active learning framework leveraging transcriptomics identifies modulators of disease phenotypes. *Science*, 390(6776):ead8577, 2025. doi: 10.1126/science.adi8577.
- [33] Rahul Satija, Jeffrey A. Farrell, David Gennert, Alexander F. Schier, and Aviv Regev. Spatial reconstruction of single-cell gene expression data. *Nature Biotechnology*, 33(5):495–502, 2015. doi: 10.1038/nbt.3192.

- 975 [34] Gábor J. Székely and Maria L. Rizzo. Energy statistics: a class of statistics based on distances. *Journal*
976 *of Statistical Planning and Inference*, 143(8):1249–1272, 2013. doi: 10.1016/j.jspi.2013.03.018.
- 977 [35] Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola.
978 A kernel two-sample test. *Journal of Machine Learning Research*, 13(25):723–773, 2012.
- 979 [36] Julien Rabin, Gabriel Peyré, Julie Delon, and Marc Bernot. Wasserstein barycenter and its applica-
980 tion to texture mixing. In *Scale Space and Variational Methods in Computer Vision (SSVM 2011)*, pages
981 435–446, 2012. doi: 10.1007/978-3-642-24785-9_37.
- 982 [37] Sayash Kapoor and Arvind Narayanan. Leakage and the reproducibility crisis in machine-learning-
983 based science. *Patterns*, 4(9):100804, 2023. doi: 10.1016/j.patter.2023.100804.
- 984 [38] Kalervo Järvelin and Jaana Kekäläinen. Cumulated gain-based evaluation of IR techniques. *ACM*
985 *Transactions on Information Systems*, 20(4):422–446, 2002. doi: 10.1145/582415.582418.
- 986 [39] Marlon Stoeckius, Christoph Hafemeister, William Stephenson, Brian Houck-Loomis, Pratip K.
987 Chattopadhyay, Harold Swerdlow, Rahul Satija, and Peter Smibert. Simultaneous epitope and tran-
988 scriptome measurement in single cells. *Nature Methods*, 14(9):865–868, 2017. doi: 10.1038/nmeth.
989 4380.
- 990 [40] Jordan W. Squair et al. Confronting false discoveries in single-cell differential expression. *Nature*
991 *Communications*, 12(1):5692, 2021. doi: 10.1038/s41467-021-25960-2.
- 992 [41] Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: a practical and powerful
993 approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1):289–300, 1995. doi:
994 10.1111/j.2517-6161.1995.tb02031.x.
- 995 [42] F. Alexander Wolf, Philipp Angerer, and Fabian J. Theis. SCANPY: large-scale single-cell gene
996 expression data analysis. *Genome Biology*, 19(1):15, 2018. doi: 10.1186/s13059-017-1382-0.
- 997 [43] Fabian Pedregosa et al. Scikit-learn: machine learning in Python. *Journal of Machine Learning Re-*
998 *search*, 12:2825–2830, 2011. JMLR registers no DOI.